

# Chapter 6

## Statistical Models for Network Graphs

Statistical Analysis of Network Data, with R - Eric D. Kolaczyk

*Thu Nguyen*

*02 July, 2019*

### Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Exponential Random Graph Models</b>	<b>2</b>
2.1	General Formulation . . . . .	2
2.2	Specifying a Model . . . . .	3
2.3	Model Fitting . . . . .	5
2.4	Goodness-of-Fit of model . . . . .	6
<b>3</b>	<b>Network Block Models</b>	<b>7</b>
3.1	Model Specification . . . . .	7
3.2	Model Fitting . . . . .	7
3.3	Goodness-of-Fit . . . . .	9
<b>4</b>	<b>Latent Network Models</b>	<b>10</b>
4.1	General Formulation . . . . .	10
4.2	Specifying the Latent Effects . . . . .	10
4.3	Model Fitting . . . . .	11
4.4	Goodness-of-Fit . . . . .	13

---

Libraries

```
library(igraph)
library(igraphdata)
library(sand)
```

---

### 1 Introduction

- **Exponential Random Graph models** ~ standard regression models, in particular *generalized linear models*
- **Stochastic Block models** ~ mixture models, of *classical random graph models*
- **Latent Network models** ~ network-based variant: using *observed + unobserved* variables in modeling

## 2 Exponential Random Graph Models

Exponential Random Graphs Models (ERGMs)  
~ Classical Generalized Linear Models (GLMs)

### 2.1 General Formulation

Let  $G = (V, E)$  be random, let  $Y_{ij} = Y_{ji} = \mathbb{1}_{\{i,j\} \in E}$ , then  $\mathbf{Y} = [Y_{ij}]$  is the random adjacency matrix. Let  $\mathbf{y} = [y_{ij}]$  be a particular graph, then:

$$P_{\theta}(\mathbf{Y} = \mathbf{y}) = \frac{1}{\kappa} \exp \left\{ \sum_H \theta_H g_H(\mathbf{y}) \right\}$$

- $H$ : configuration, defined to be a set of possible edges among a subset of vertices in  $G$
- $g_H(\mathbf{y}) = \prod_{y_{ij} \in H} y_{ij}$ : an indicator of  $H$ : 1 if  $H$  occurs in  $\mathbf{y}$ , and 0 otherwise
- $\theta_H \neq 0 \iff Y_{ij}$  are dependent for all pairs of vertices in  $H$
- $\kappa$ : normalization constant

Recall that a random vector  $\mathbf{Z}$  belongs to an *exponential family* if the pdf is:

$$P_{\theta}(\mathbf{Z} = \mathbf{z}) = \exp \{ \theta^T \mathbf{g}(\mathbf{z}) - \psi(\theta) \}$$

where  $\theta$ :  $p \times 1$  vector of parameters,  $\mathbf{g}(\cdot)$ :  $p$  dimensional function of  $\mathbf{z}$ , and  $\psi(\theta)$ : normalization constant.

To draw *ERGMs* in R, use [ergm](#) package, part of the [statnet](#) package suite.

```
data(lazega)
A <- get.adjacency(lazega) # 1: create adjacency matrix
v.attrs <- get.data.frame(lazega, what = 'vertices') # 2: arrange into data frame
library(ergm)
lazega.s <- network::as.network(as.matrix(A), # 3: network object for ergm
                               directed = FALSE)
network::set.vertex.attribute(lazega.s, "Office", v.attrs$Office)
network::set.vertex.attribute(lazega.s, "Practice", v.attrs$Practice)
network::set.vertex.attribute(lazega.s, "Gender", v.attrs$Gender)
network::set.vertex.attribute(lazega.s, "Seniority", v.attrs$Seniority)
```

## 2.2 Specifying a Model

Given a general formulation of *ERGM*, we can introduce some assumptions to get different models:

1. *Bernoulli random graph*: the probability of any edge between every pair of vertices is **iid**, giving

$$P_{\theta}(\mathbf{Y} = \mathbf{y}) = \frac{1}{\kappa} \exp \left\{ \sum_{i,j} \theta_{ij} y_{ij} \right\}$$

2. *Homogeneity*: given the Bernoulli assumption, let  $\theta_{ij} = c$  for all  $i, j$ , where  $c$  is some constant, giving

$$P_{\theta}(\mathbf{Y} = \mathbf{y}) = \frac{1}{\kappa} \exp \left\{ \theta L(\mathbf{y}) \right\}$$

where  $L(\mathbf{y}) = |E_G|$ . Note that this is equivalent to a *Bernoulli random graph* with  $p = \frac{\exp(\theta)}{1+\exp(\theta)}$ .

---

In R: to specify the model, use `formula()`:

```
my.ergm.bern <- formula(lazega.s ~ edges)
summary(my.ergm.bern)
```

```
## edges
##      115
```

Furthermore, suppose that we want to incorporate statistics of higher-order global network structure such as  $k$ -stars  $S_k(\mathbf{y})$ , and triangles  $T(\mathbf{y})$ , ... In R: in `formula()`, specify `kstar()` and `triangle`:

```
my.ergm <- formula(lazega.s ~ edges + kstar(2) + kstar(3) + triangle)
summary(my.ergm)
```

```
##      edges      kstar2      kstar3 triangle
##      115         926       2681       120
```

Additional statistics such as *Alternating k-star statistic* `altkstar()`, *Geometrically weighted degree count* `gwdegree()`, and *generalization of Triadic structures* `gwesp()`:

```
my.ergm <- formula(lazega.s ~ edges + altkstar(1, fixed = TRUE))
ergm.altkstar <- summary(my.ergm)[2]
my.ergm <- formula(lazega.s ~ edges + gwdegree(1, fixed = TRUE))
ergm.gwdegree <- summary(my.ergm)[2]
my.ergm <- formula(lazega.s ~ edges + gwesp(1, fixed = TRUE))
ergm.gwesp <- summary(my.ergm)[2]
data.frame(Statistics = c('Alternating k-star', 'Geom. weighted Degree', 'Triadic Structure'),
           Values = c(ergm.altkstar, ergm.gwdegree, ergm.gwesp))
```

```
##              Statistics      Values
## altkstar.1      Alternating k-star 196.0000
## gwdeg.fixed.1 Geom. weighted Degree  79.2000
## gwesp.fixed.1   Triadic Structure 213.1753
```

To measure the total similarity among the vertices in a network, we can look at the statistics:

$$g(\mathbf{y}, \mathbf{x}) = \sum_{1 \leq i \leq j \leq n} y_{ij} h(\mathbf{x}_i, \mathbf{x}_j)$$

where  $h$  is a symmetric function of choice, and  $\mathbf{x}_i$  is a vector of observed attributes:

1. *Main effects*:  $h(x_i, x_j) = x_i + x_j$ , in ‘R’: `nodemain()`
2. *Second-order/Homophily effects*:  $h(x_i, x_j) = \mathbb{1}_{x_i = x_j}$ , in ‘R’: `nodematch()`

```
lazega.ergm <- formula(lazega.s ~ edges
+ gwesp(log(3), fixed = TRUE)
+ nodemain("Seniority")
+ nodemain("Practice")
+ match("Practice")
+ match("Gender")
+ match("Office"))
```

---

## 2.3 Model Fitting

In general, given *iid* realizations, *ERGMs* are fit using *Maximum Likelihood Estimator*. In R, `ergm()`:

```
set.seed(42)
lazega.ergm.fit <- ergm(lazega.ergm)
anova(lazega.ergm.fit)

summary(lazega.ergm.fit)

##
## =====
## Summary of model fit
## =====
##
## Formula:   lazega.s ~ edges + gwesp(log(3), fixed = TRUE) + nodemain("Seniority") +
##            nodemain("Practice") + match("Practice") + match("Gender") +
##            match("Office")
##
## Iterations: 2 out of 20
##
## Monte Carlo MLE Results:
##
##            Estimate Std. Error MCMC % z value Pr(>|z|)
## edges          -7.00655    0.67114      0 -10.440 < 0.0001 ***
## gwesp.fixed.1.09861228866811  0.59166    0.08554      0  6.917 < 0.0001 ***
## nodecov.Seniority      0.02456    0.00620      0  3.962 < 0.0001 ***
## nodecov.Practice      0.39455    0.10218      0  3.861 0.000113 ***
## nodematch.Practice     0.76966    0.19060      0  4.038 < 0.0001 ***
## nodematch.Gender      0.73767    0.24362      0  3.028 0.002463 **
## nodematch.Office      1.16439    0.18753      0  6.209 < 0.0001 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##      Null Deviance: 873.4  on 630  degrees of freedom
## Residual Deviance: 459.6  on 623  degrees of freedom
##
## AIC: 473.6    BIC: 504.7    (Smaller is better.)
```

### Interpretation:


- practicing corporate law, not litigation, increases cooperation by  $\exp(.39455) \approx 1.48$ , or nearly 50%
- being of the same gender more than doubles the odds:  $\exp(.73767) = 2.09$ .

## 2.4 Goodness-of-Fit of model


To access GOF as fit by `ergm`, in R: `gof()` runs the Monte Carlo simulation:

```
gof.lazega.ergm <- gof(lazega.ergm.fit)
```

```
par(mfrow=c(1,3))  
plot(gof.lazega.ergm)
```



## Goodness-of-fit diagnostics



### 3 Network Block Models

Network Block Models  $\sim$  Classical Mixture Models

#### 3.1 Model Specification

Let  $G = (V, E)$ , with adjacency matrix  $\mathbf{Y} = [Y_{ij}]$ . Suppose each vertex  $i \in V_G$  belongs to 1 of  $Q$  classes  $\mathcal{C}_1, \dots, \mathcal{C}_Q$ , and the class label is known:  $q = q(i), \forall i$ . Conditioned on class labels  $q, r$  or vertices  $i, j$ , a *block model* is such that each  $Y_{ij}$  is *iid Bernoulli* with probability  $\pi_{ij}$ . For undirected graph:  $\pi_{ij} = \pi_{ji}$ , giving

$$P_{\theta}(\mathbf{Y} = \mathbf{y}) = \frac{1}{\kappa} \exp \left\{ \sum_{q,r} \theta_{qr} L_{qr}(\mathbf{y}) \right\}$$

If given  $Q$  classes but *unknown* class labels, the model becomes *Stochastic Block Model (SBM)*.

#### 3.2 Model Fitting

In a *non-stochastic block model*, the edge probabilities  $\pi_{qr}$  are estimated using *Maximum Likelihood Estimates*. In R: `mixer()` from the `mixer` package, to specify  $\min(Q)$  `qmin`,  $\max(Q)$  `qmax`:

```
library(mixer)
set.seed(42)
fblog.sbm <- mixer(as.matrix(get.adjacency(fblog)), qmin = 2, qmax = 15)
```

## Mixer: the adjacency matrix has been transformed in a undirected edge list

```
fblog.sbm.output <- getModel(fblog.sbm)
names(fblog.sbm.output)
```

```
## [1] "q"          "criterion" "alphas"    "Pis"       "Taus"
```

The criterion above is *Integration Classification Likelihood (ICL)*, similar to AIC and BIC.

```
print(paste0('Fitted model: q = ', fblog.sbm.output$q))
```

```
## [1] "Fitted model: q = 12"
```

```
print(cat('Estimated proportions:', '\n', round(fblogger.sbm.output$alphas, 3), '\n'))
```

```
## Estimated proportions:
##  0.151 0.127 0.108 0.057 0.136 0.031 0.124 0.093 0.01 0.021 0.125 0.016
## NULL
```

Thus, *Stochastic Block models* can be used for *graph partitioning*.

```
round(fblog.sbm.output$Taus[, 1:3], 6)
```

```
##           [,1]      [,2]      [,3]
## [1,] 0.999975 0.004911 0.999981
## [2,] 0.000014 0.000000 0.000001
## [3,] 0.000010 0.994221 0.000018
## [4,] 0.000000 0.000000 0.000000
## [5,] 0.000001 0.000000 0.000000
## [6,] 0.000000 0.000000 0.000000
## [7,] 0.000000 0.000868 0.000000
## [8,] 0.000000 0.000000 0.000000
## [9,] 0.000000 0.000000 0.000000
## [10,] 0.000000 0.000000 0.000000
## [11,] 0.000000 0.000000 0.000000
## [12,] 0.000000 0.000000 0.000000
```

**Interpretation:**

- $P(v_1 \in \mathcal{C}_1) = 0.9999747$ , or the model labels vertex 1 in class 1
- $P(v_2 \in \mathcal{C}_3) = 0.9942208$ , and so on

---

*Entropy* of a discrete pmf  $\mathbf{p} = (p_1, \dots, p_Q) : H(\mathbf{p}) := -\sum_{q=1}^Q p_q \log_2(p_q)$ : smaller  $H(\mathbf{p})$  indicates distribution is concentrated on fewer classes.

```
my.ent <- function(x) { -sum(x * log(x, 2)) }      # fn to calculate entropy
apply(fblog.sbm.output$Taus[, 1:3], 2, my.ent)
```

```
## [1] 0.0004559176 0.0548083843 0.0003231649
```

```
print(paste('Entropy for if classes are Uniformly dist.:', log(fblog.sbm.output$q, 2)))
```

```
## [1] "Entropy for if classes are Uniformly dist.: 3.58496250072116"
```

**Interpretation:** small entropy values are consistent with the fblog network: vertices are concentrated on few classes.



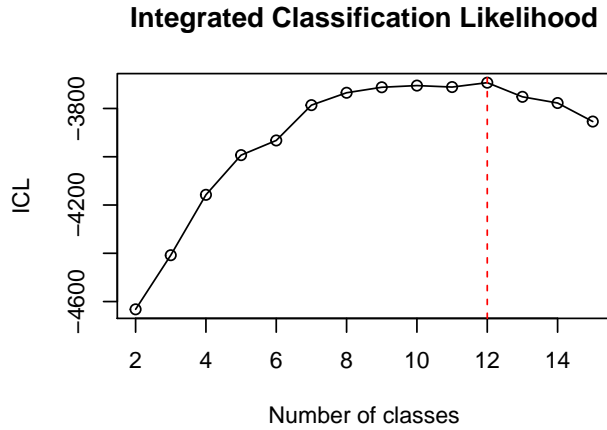
```
summary(apply(fblog.sbm.output$Taus, 2, my.ent))
```

```
##      Min.    1st Qu.    Median      Mean   3rd Qu.      Max.
## 0.0000000 0.0000000 0.0000006 0.0414403 0.0029971 1.4130217
```




### 3.3 Goodness-of-Fit

```
plot(fblog.sbm, classes=as.factor(V(fblog)$PolParty))
```



**Inter/intra class probabilities**



#### Interpretation:

1. while  $Q = 12$  gives  $\max(ICL)$ ,  $Q \in \{8, 9, 10, 11, 12\}$  are also reasonable choices
2. from adjacency matrix  $\mathbf{Y}$ , there are 5 smaller classes and 7 larger classes, whose vertices tend to primarily connect within classes, and with vertices from only some certain other classes
3. degree distribution: blue curve: fitted *SBN*, vs. yellow: observed distribution.

## 4 Latent Network Models

---

### 4.1 General Formulation

Given the absence of any co-variate information, it is natural to assume exchangeability of vertices in  $G = (V, E)$ , with adjacency matrix  $\mathbf{Y} = [Y_{ij}]$ , giving

$$Y_{ij} = h(\alpha, u_i, u_j, \epsilon_{ij}) \quad (1)$$

where  $\alpha$ : constant,  $u_i$ : *iid* latent variables, and  $\epsilon_{ij}$ : *iid* pair-specific effects, and  $h$ : symmetric wrt  $u_i, u_j$ . An example is *probit model*:

$$\mathbb{P}(\mathbf{Y}_{ij} = 1 \mid \mathbf{X}_{ij} = \mathbf{x}_{ij}) = \phi(\alpha + \mathbf{x}_{ij}^T \beta + \alpha(u_i, u_j)) \quad (2)$$

where  $\phi$  is CMF of  $\mathbf{Z} \sim \mathcal{N}(0, 1)$ . Let  $p_{ij} = (2)$ , the conditional model for  $\mathbf{Y}$  is

$$\mathbb{P}(\mathbf{Y} = \mathbf{y} \mid \mathbf{X}, u_1, \dots, u_n) = \prod_{i < j} p_{ij}^{y_{ij}} (1 - p_{ij})^{1 - y_{ij}}$$


---

### 4.2 Specifying the Latent Effects

From (1), the function  $\alpha(\cdot, \cdot)$  dictates the effects of the latent variables, in particular:

1. *Latent class models*: analogous to *Stochastic block models* above:  $u_i \in \{1, \dots, Q\}$ ,  $\alpha(u_i, u_j) = m_{u_i u_j}$  symmetrically
  2. *Latent distance model*: under the *principle of homophily*: vertices with more similar characteristics tend to establish an edge,  $\alpha(u_i, u_j) = -|u_i - u_j|$  for some distance metric
  3. *Eigenmodel*: under the *principles of eigen-analysis*:  $\alpha(u_i, u_j) = a_i^T \Lambda u_j$ , where  $u_i$ :  $Q$ -length random vectors, and  $\Lambda$ :  $Q \times Q$  diagonal matrix;
    - note that if  $\mathbf{U} = [u_1, \dots, u_Q]$ , then  $\mathbf{U} \Lambda \mathbf{U}^T$  is analogous to eigen-decomposition of all pairwise latent effects  $\alpha(u_i, u_j)$
    - *Eigenmodels* can be thought of as a generalization of both *Latent class* and *Latent distance* models.
-

### 4.3 Model Fitting

Given the `lazega` dataset, it is natural to hypothesize that collaboration is driven by:

- similarity of practice  $\sim$  a form of *homophily*, or
- similarity of office location  $\sim$  a proxy for *distance*

As such, we can compare 3 fitted models with different settings:

1. no pair-specific covariates
2. a covariate for common practice
3. a covariate for shared office location

In R: `eigenmodel` package runs the Monte Carlo Markov Chain simulation to obtain the posterior distributions from the conjugate priors (*Bayesian* approach).

Model 1: no pair-specific covariates:

```
library(eigenmodel)
set.seed(42)
A <- get.adjacency(lazega, sparse = FALSE)
lazega.leig.fit1 <- eigenmodel_mcmc(A, R = 2, S = 11000, burn = 10000)
```

Model 2: a covariate for common practice:

```
# Common practice effects
same.prac.op <- v.attr.lazega$Practice %o% v.attr.lazega$Practice
same.prac <- matrix(as.numeric(same.prac.op %in% c(1,4,9)), 36, 36)
same.prac <- array(same.prac, dim = c(36, 36, 1))
# Fit model
lazega.leig.fit2 <- eigenmodel_mcmc(A, same.prac, R = 2, S = 11000, burn = 10000)
```

Model 3: a covariate for shared office location

```
# Common office effects
same.off.op <- v.attr.lazega$Office %o% v.attr.lazega$Office
same.off <- matrix(as.numeric(same.off.op %in% c(1,4,9)), 36, 36)
same.off <- array(same.off, dim = c(36, 36, 1))
# Fit model
lazega.leig.fit3 <- eigenmodel_mcmc(A, same.off, R = 2, S = 11000, burn = 10000)
```


To compare the representation, we extract and plot the eigenvectors for each model, using `eigen()`:

```
lat.sp.1 <- eigen(lazega.leig.fit1$ULU_postmean)$vec[, 1:2]
lat.sp.2 <- eigen(lazega.leig.fit2$ULU_postmean)$vec[, 1:2]
lat.sp.3 <- eigen(lazega.leig.fit3$ULU_postmean)$vec[, 1:2]


par(mfrow=c(1,3)); par(mar=c(0,1,1,1))
v.colors <- c('red', 'blue', 'yellow')[V(lazega)$Office]
v.shapes <- c('circle', 'square')[V(lazega)$Practice]
v.size <- 3.5*sqrt(V(lazega)$Years)
v.label <- V(lazega)$Seniority
plot(lazega, layout = lat.sp.1, vertex.color = v.colors, vertex.shape = v.shapes,
     vertex.size = v.size, vertex.label = v.label, main = 'Model 1: No covariates')
```

```
plot(lazega, layout = lat.sp.2, vertex.color = v.colors, vertex.shape = v.shapes,
     vertex.size = v.size, vertex.label = v.label, main = 'Model 2: Common practice')
plot(lazega, layout = lat.sp.3, vertex.color = v.colors, vertex.shape = v.shapes,
     vertex.size = v.size, vertex.label = v.label, main = 'Model 3: Shared office')
```


Model 1: No covariates



Model 2: Common practice



Model 3: Shared office



### Interpretation:

- Models 1 and 2: lawyers are clustered into 2 main groups based on office location
- Model 3: common practice appears to not well distinguish the network structure

```
lambda1 <- apply(lazega.leig.fit1$L_postsamp, 2, mean)
lambda2 <- apply(lazega.leig.fit2$L_postsamp, 2, mean)
lambda3 <- apply(lazega.leig.fit3$L_postsamp, 2, mean)
data.frame(Models = c('No covariate', 'Common practice', 'Common office'),
           lambda.1 = c(lambda1[1], lambda2[1], lambda3[1]),
           lambda.2 = c(lambda1[2], lambda2[2], lambda3[2]))
```

```
##           Models  lambda.1  lambda.2
## 1   No covariate 0.2876963 0.9709850
## 2 Common practice 0.9085345 -0.1531384
## 3   Common office 0.4031347 0.1892327
```

The table is consistent with the plot: in Model 1 and 2, there is 1 value eigenvalue  $\lambda$  that dominates, while in model 3, the 2 eigenvalues do not differ much.

## 4.4 Goodness-of-Fit

Here, we use cross-validation, with  $k = 5$ , to assess GOF for the 3 different models.

```
perm.index <- sample(1:630)
nfolds <- 5
nmiss <- 630/nfolds
Avec <- A[lower.tri(A)]
Avec.pred1 <- numeric(length(Avec))
Avec.pred2 <- numeric(length(Avec))
Avec.pred3 <- numeric(length(Avec))

for (i in seq(1, nfolds)) {
  # Index of missin values
  miss.index <- seq(((i-1)*nmiss + 1), i*nmiss, 1)
  A.miss.index <- perm.index[miss.index]

  # Fill a new Atemp with NAs
  Avec.temp <- Avec
  Avec.temp[A.miss.index] <- rep('NA', length(A.miss.index))
  Avec.temp <- as.numeric(Avec.temp)
  Atemp <- matrix(0, 36, 36)
  Atemp[lower.tri(Atemp)] <- Avec.temp
  Atemp <- Atemp + t(Atemp)

  # Fit model and predict, model 1
  Y <- Atemp
  model1.fit <- eigenmodel_mcmc(Y, R = 2, S = 11000, burn = 10000)
  model1.pred <- model1.fit$Y_postmean
  model1.pred.vec <- model1.pred[lower.tri(model1.pred)]
  Avec.pred1[A.miss.index] <- model1.pred.vec[A.miss.index]

  # Fit model and predict, model 2
  model2.fit <- eigenmodel_mcmc(Y, same.prac, R = 2, S = 11000, burn = 10000)
  model2.pred <- model2.fit$Y_postmean
  model2.pred.vec <- model2.pred[lower.tri(model2.pred)]
  Avec.pred2[A.miss.index] <- model2.pred.vec[A.miss.index]

  # Fit model and predict, model 3
  model3.fit <- eigenmodel_mcmc(Y, same.off, R = 2, S = 11000, burn = 10000)
  model3.pred <- model3.fit$Y_postmean
  model3.pred.vec <- model3.pred[lower.tri(model3.pred)]
  Avec.pred3[A.miss.index] <- model3.pred.vec[A.miss.index]
}
```


Similar to *Classification* problem, 1 way to evaluate the fitted model is through the ROC curve and the AUC percentage from the `ROCR` package: `prediction()` and `performance()`:

```
par(mar=c(5,5,1,2)); par(oma=c(0, 0, 0, 5))
library(ROCR)
pred1 <- prediction(Avec.pred1, Avec)
perf1 <- performance(pred1, 'tpr', 'fpr')
pred2 <- prediction(Avec.pred2, Avec)
perf2 <- performance(pred2, 'tpr', 'fpr')
pred3 <- prediction(Avec.pred3, Avec)
perf3 <- performance(pred3, 'tpr', 'fpr')
```

```

plot(perf1, col = 'blue', lwd = 3, legend = 'Model 1')
plot(perf2, col = 'red', lwd = 3, add = TRUE, legend = 'Model 2')
plot(perf3, col = 'brown', lwd = 3, add = TRUE)
legend(par('usr')[2], par('usr')[4], xpd=NA, bty = 'n',
       legend = c('Model 1', 'Model 2', 'Model 3'),
       col = c('blue', 'red', 'brown'), c('blue', 'red', 'brown'))

```



```

auc.mod1 <- slot(performance(pred1, 'auc'), 'y.values')[[1]]
auc.mod2 <- slot(performance(pred2, 'auc'), 'y.values')[[1]]
auc.mod3 <- slot(performance(pred3, 'auc'), 'y.values')[[1]]
data.frame(Models = c('No covariate', 'Common practice', 'Common office'),
           AUC = c(auc.mod1, auc.mod2, auc.mod3))

```

```

##           Models      AUC
## 1   No covariate 0.8172309
## 2 Common practice 0.8324187
## 3   Common office 0.8284255

```

**Comment:** all models appear to be comparable in their performance and to perform well: AUC of over 80%.